

Machine Learning Assignment 1

Kyriacos Xanthos
CID: 01389741

February 3, 2022

Question 1

In this question we are given a dataset of the form $\mathcal{D} = \{x_i, z_i, y_i\}_{i=1}^{100}$ and we consider a model of the form:

$$Y_i = aZ_i + (b_1 + b_2Z_i) \cos(X_i) + (c_1 + c_2Z_i) X_i + (d_1 + d_2Z_i) X_i^2 + (e_1 + e_2Z_i) X_i^3 + \epsilon_i \quad (1)$$

where ϵ_i are independent random variables with mean 0.

1.

In this part we consider ridge regression. Note that ridge regression is an extension to linear regression by introducing an l_2 penalty on the coefficients to a standard least squares problem. Mathematically we are trying to solve the optimization problem:

$$\underset{\beta}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda (\beta^T \beta) \quad (2)$$

where in our example $\beta = [a \ b_1 \ b_2 \ c_1 \ c_2 \ d_1 \ d_2 \ e_1 \ e_2]^T$, λ is the regularization parameter, and \mathbf{y} and \mathbf{X} are defined as:

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}, \quad \mathbf{X}^T = \begin{bmatrix} Z_1 & Z_2 & \dots & Z_n \\ \cos(X_1) & \cos(X_2) & \dots & \cos(X_n) \\ Z_1 \cos(X_1) & Z_2 \cos(X_2) & \dots & Z_n \cos(X_n) \\ X_1 & X_2 & \dots & X_n \\ Z_1 X_1 & Z_2 X_2 & \dots & Z_n X_n \\ X_1^2 & X_2^2 & \dots & X_n^2 \\ Z_1 X_1^2 & Z_2 X_2^2 & \dots & Z_n X_n^2 \\ X_1^3 & X_2^3 & \dots & X_n^3 \\ Z_1 X_1^3 & Z_2 X_2^3 & \dots & Z_n X_n^3 \end{bmatrix}$$

where n is 100 for our dataset.

We will use the function **Ridge** from the package **sklearn** in python which solves equation 2 and estimates the coefficients in the following form:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y} \quad (3)$$

We set the **fit_intercept** parameter to **False** so that it does not include an intercept and use the **svd** solver which performs singular value Decomposition at a low computational cost. We will assess our model performance by using the root mean squared error, defined as:

$$\text{rmse} = \left(\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \right)^{1/2} \quad (4)$$

where $\hat{y} = \mathbf{X}\beta$. With `sklearn` this is calculated using the function `mean_squared_error` with the argument `squared=False`.

λ is actually a hyperparameter and we need to tune it so that we get the optimum value. To tune it we use cross validation. The idea is that we split the data \mathcal{D} randomly, into K -folds. We then fit the model K times, we validate on k^{th} fold and train on the remaining $K - 1$ folds. By averaging across multiple folds of the data we decrease the variability in the model fit and validation error caused by the randomness of the splits. Figure 1 shows how the model RMSE varies for different parameters λ and for $K = 5$ folds. It is clear that many of the folds have not reached a minimum and suggest that the minimum is less than 0. However, we cannot have $\lambda < 0$ because then the matrix $(\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I})^{-1}$ would not be invertible.




Figure 1: RMSE as defined in 4 against different values of hyperparameter λ for different folds.

We utilise cross validation using the `GridSearchCV` function from `sklearn` to search for the optimum value of λ that minimises the RMSE. This value is found to be $\lambda = 1.2156$ and looking at figure 2 we see that it indeed minimises the RMSE. We interpret this result with a grain of salt since the splits in figure 1 did not show a clear minimum before zero but this might be because our dataset only has 100 values and so the splits just consist of 20 points.




Figure 2: RMSE as defined in 4 against different values of hyperparameter λ .

Using this optimum value of λ we can re-run the Ridge regression and obtain the model summary in figure 3. The model seems to predict the true y with a high accuracy, with $\text{RMSE} = 0.286$ and the residual plot shows that there is not any particular bias. The coefficients are found to be:

$$\hat{\beta} = [0.0159 \quad 0.9993 \quad 0.8435 \quad 0.0242 \quad 0.01536 \quad -0.0006 \quad 0.1834 \quad 0.01453 \quad -0.0006]^T$$




Figure 3: Summary of Ridge regression model with $\lambda = 1.2156$. Comparison of predicted y values with true y values (left) and residual plot (right)

Now we try to predict the value of Y^* for $X^* \in \{-7, -6.9, -6.8, \dots, 6.8, 6.9, 7\}$ and $Z^* \in \{0, 1\}$. The equation we use to predict this is $Y^* = X^* \hat{\beta}$ where $\hat{\beta}$ are the model parameter from our fit above.

We can see the fit of these two models in figure 4. It is clear that our model performs well for both z values. The fit seems to follow the trend of all points and even captures the few points which have higher values of y . This confirms a good model qualitatively on unseen data.




Figure 4: Fit of the model using Ridge regression on unseen data for $z = 0$ (left) and $z = 1$ (right).

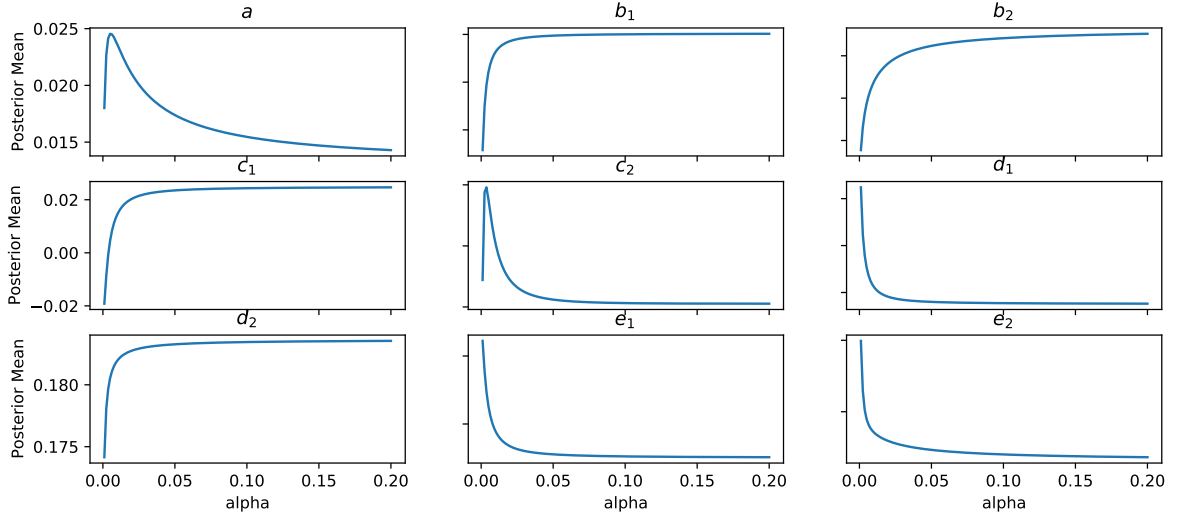


Figure 5: Posterior mean of 9 model parameters as a function of α

2.

We now perform Bayesian inference to determine the same parameters of this problem. The fundamental idea is that these parameters cannot be considered as single values, instead they acquire probability distributions. We are therefore interested in the posterior distribution over the parameters β ,

$$p(\beta | \mathbf{X}, \mathbf{y}, \sigma) = \frac{p(\mathbf{y} | \beta, \mathbf{X}, \sigma)p(\beta)}{p(\mathbf{y} | \mathbf{X}, \sigma)} \quad (5)$$

where $\sigma^2 = 0.01$ and the marginal likelihood is written as $p(\mathbf{y} | \mathbf{X}, \sigma) = \int p(\mathbf{y} | \beta, \mathbf{X}, \sigma)p(\beta)d\beta$. We will consider a Gaussian prior of the form $p(\beta) = \mathcal{N}_\beta(\mathbf{0}, \Lambda)$, where $\Lambda = \alpha \mathbf{I}$, the multivariate normal centred at zero and with scaled identity covariance which encodes the belief that the parameters should be small which means that α is a parameter which we choose.

This is solved by [1]

$$\begin{aligned} p(\beta | \mathbf{X}, \mathbf{y}, \sigma) &= \frac{\mathcal{N}_\mathbf{y}(\mathbf{x}\beta, \sigma\mathbf{I})\mathcal{N}_\beta(\mathbf{0}, \Lambda)}{\int \mathcal{N}_\mathbf{y}(\mathbf{X}\beta, \sigma\mathbf{I})\mathcal{N}_\beta(\mathbf{0}, \Lambda)d\beta} \\ &= \mathcal{N}_\beta(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \end{aligned} \quad (6)$$

where

$$\boldsymbol{\mu} = \left(\mathbf{X}^T \mathbf{X} + \frac{\sigma^2}{\alpha} \mathbf{I} \right)^{-1} \mathbf{X}^T \mathbf{y} \quad , \quad \boldsymbol{\Sigma} = \sigma^2 \left(\mathbf{X}^T \mathbf{X} + \frac{\sigma^2}{\alpha} \mathbf{I} \right)^{-1} \quad (7)$$

Calculating 6 in `python` we can find the probability distribution of each parameter. Plotting the posterior mean of the 9 model parameters as a function of α we obtain figure 5.

It is clear that as the value of α increases, all of the 9 parameters converge either to a maximum or a minimum value. This means that for any $\alpha > 0.2$ we can be confident that the posterior mean of the parameter is converged.

To calculate the predictive posterior distribution we use

$$\begin{aligned} p(\mathbf{y}_{\text{new}} | \mathbf{x}_{\text{new}}, \mathbf{X}, \mathbf{y}, \sigma) &= \int p(\mathbf{y}_{\text{new}} | \mathbf{x}_{\text{new}}, \beta, \sigma) p(\beta | \mathbf{X}, \mathbf{y}, \sigma) d\beta \\ &= \mathcal{N}_{\mathbf{y}_{\text{new}}} \left(\mathbf{x}_{\text{new}}^T \boldsymbol{\mu}, \mathbf{x}_{\text{new}}^T \boldsymbol{\Sigma} \mathbf{x}_{\text{new}} + \sigma^2 \right) \end{aligned} \quad (8)$$

as stated in the notes [1]. We again use the same new values as above, namely $X^* \in \{-7, -6.9, -6.8, \dots, 6.8, 6.9, 7\}$ and $Z^* \in \{0, 1\}$ to obtain the new matrix \mathbf{x}_{new} . We then calculate the posterior mean and its 95% credible interval. Credible intervals are analogous to the confidence intervals in frequentist statistics but for the bayesian interpretation. They determine the distribution of possible values of the posterior mean [2]. We calculate the




Figure 6: Mean of Posterior Predictive Distribution as a function of X with $z = 0$ (left) and $z = 1$ (right). The shaded blue shows the 95% credible interval of the computation.

credible interval using the 2.5% and 97.5% quantiles of the Gaussian distribution with the parameters stated in equation 6. The plots for $z = 0$ and $z = 1$ are shown in figure 6.

The fit of the model seems to be very good for both values of z . The 95% credible interval shows that it covers most of the points in our data except some of the points that exhibit local maxima and local minima. If the range of the credible interval was increased to 90% then the interval would cover all points.

3.

Now we will make predictions on $(X, Z) \in \{-5, 7\} \times \{0, 1\}$. We will use 3 different values of α for our predictions. The first one consists of understanding the the regularised ridge least squares problem is exactly the maximum a posteriori estimate of the parameters β with the regularisation parameter $\lambda = \sigma^2/\alpha$. So then using the optimum λ found in part 1, $\alpha = 0.1/1.2156 = 0.0823$. The second α will be a value much smaller which would mean that the parameters in figure 5 would not have converged (here we use $\alpha = 0.001$) and the third value would be a much larger value that ensures that the posterior mean of the parameters has converged ($\alpha = 1$).

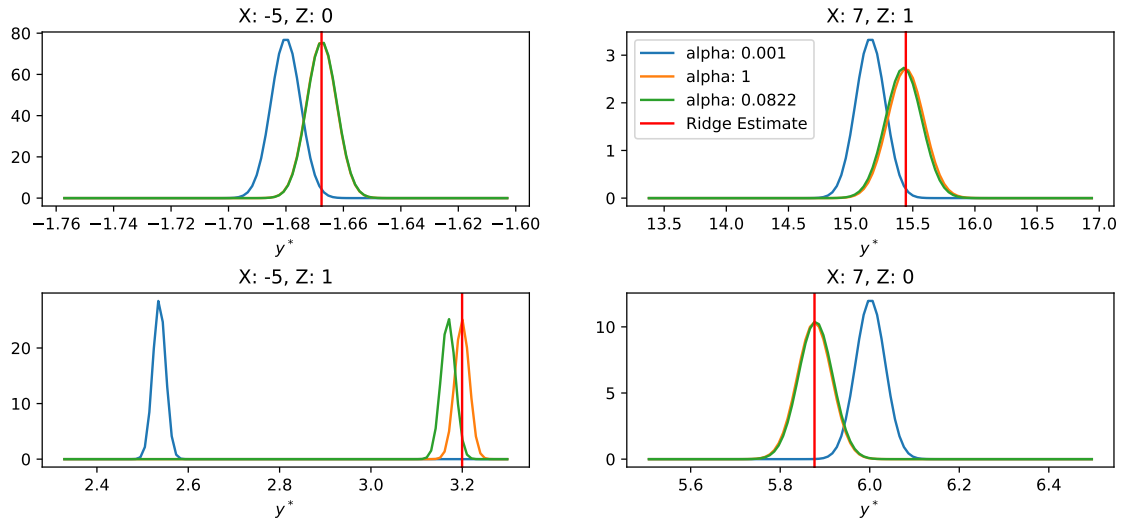


Figure 7: PDF of Predictive distributions for different observation points.

Looking at figure 7 we can observe that the posterior distributions for $\alpha = 0.0822, 1$ are very close for all 4 observations and this makes sense because as observed in figure 5 our parameters seemed to almost converge after 0.1. For $\alpha = 0.001$ as expected in all 4

observations the posterior distribution is far away from the other 2 values especially for $X = -5, Z = 1$ which shows that the prediction would not be accurate. Indeed, using the relationship with the λ of the ridge regression we see that a very small value of α means there is heavy regularisation in the problem which was not the case when we searched for the optimum λ using cross validation. The red vertical line shows that the predictions using Bayesian inference are almost equivalent to the Ridge regression problem, as expected.

Question 2

We now consider a slightly modified version of the Million Song Dataset [3]. For each song in the dataset there are 91 attributes with the first one being the year it was published. We will create models that predict if a song was produced before or after 1980. This means we will change the first attribute to a binary variable for each song, namely it will be a 0 if it was produced before 1980 and 1 if it was produced after. We can see the distribution of songs for each year in figure 8.




Figure 8: Distribution of songs produced each year. Grey are all the songs before 1980 and blue all the songs after.

We can observe that the number of songs after 1955 has been right censored to 200 counts. This was done in order to have roughly the same number of songs before and after 1980. Indeed when we count them, we have a dataset with 6201 songs after 1980 and 6322 songs before, a balanced dataset.

Since we have 90 attributes, we cannot visualise the attributes to explore some of their features. We note that we have no missing data, and therefore we can right away use classification. In order to be able to assess the efficiency of our classifiers, we split the data to training and testing using the function `train_test_split` from `sklearn` which shuffles the data and then randomly splits it in training and testing subsets. Since we have a large dataset we use a test size of 30% to reduce the computational cost in training. We then use the `StandardScaler` function of `sklearn` to scale all of the attributes of the songs so that they have mean 0 and variance 1. We do that so that the effect of each attribute will be captured in our model and we make sure to do the scaling after the split of train and testing sets.

LDA

The Linear Discriminant Analysis (LDA) classifier partitions our data into regions with the same class predictions via separating hyperplanes. This is a generative approach that models the class conditional densities $p(x | y = k)$ and class priors $p(y = k)$. Then using Bayes Theorem we compute the posterior probabilities

$$p(y = k | x) = \frac{p(x | y = k)p(y = k)}{p(x)} \quad (9)$$

where $p(x) = \sum_{j=1}^K p(x | y = j)p(y = j)$ [1].

Then it assumes that the class-conditional densities are Gaussian,

$$x | y = k \sim \mathcal{N}(\mu_k, \Sigma), \quad \text{for } 1 \leq k \leq K \quad (10)$$

where Σ is the covariance matrix and μ_k are the means. Then to create the boundaries we use parameter estimation using Maximum Likelihood. These estimates are

$$\begin{aligned} \hat{\pi}_k &= \frac{N_k}{N} \\ \hat{\mu}_k &= \frac{1}{N_k} \sum_{i, y^{(i)}=k} x^{(i)} \\ \hat{\Sigma} &= \frac{1}{N} \sum_{k=1}^K \sum_{i, y^{(i)}=k} (x^{(i)} - \mu_k) (x^{(i)} - \mu_k)^T \end{aligned} \quad (11)$$

where N_k are the number of i s such that $y^{(i)} = k$, N are the total number of points.

To calculate this we use the `LinearDiscriminantAnalysis` function of `sklearn`. It has a linear decision boundary which is generated by fitting Gaussian class conditional densities to the data assuming that all classes share the same covariance matrix [4]. We assume no shrinkage, and use the `svd` solver which performs singular value decomposition. We could also use the `lsqr` method which uses the dedicated least squares routine and it is the fastest solver but it produced less accurate results so we preferred `svd`.

We then fit this solver to our training data and make predictions on our testing data. A common metric for the performance of our classifier is a confusion matrix [5] which summarises the correct and false classifications.

| | | |
|-------|-------|------|
| | FALSE | TRUE |
| FALSE | 1470 | 428 |
| TRUE | 475 | 1384 |

| | | |
|-------|-------|------|
| | FALSE | TRUE |
| FALSE | 1478 | 459 |
| TRUE | 479 | 1341 |

| | | |
|-------|-------|------|
| | FALSE | TRUE |
| FALSE | 994 | 910 |
| TRUE | 355 | 1498 |

Table 1: Confusion Matrix of Test Set (30%) using LDA (left), KNN (middle) and Naive Bayes (right)

As seen in table 1 our classifier was able to classify 76% of the songs correctly if they were before or after 1980 (Accuracy) with error rate 24%. Another useful metric for the performance of classifiers are the Receiver Operatic Characteristic Curves (ROC) [6]. They illustrate the diagnostic ability of a binary classifier as its discrimination threshold is varied. We plot the True Positive Rate

$$\text{TPR} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

against the False Negative Rate ($\text{FNR} = 1 - \text{TPR}$)¹. The perfect classifier would exhibit a vertical line to $\text{TPR} = 1$ and then a horizontal line to $\text{FPR} = 1$. We can see the ROC Curve of our classifier in figure 10 and it clear that it follows the pattern we expect with the curve peaking relatively close to to the (0,1) point.

A quantitative metric that can judge the performance of our classifier is the AUC (Area Under the Curve) of an ROC curve. This area under the curve is equal to the probability that the classifier will rank a randomly chosen positive instance higher than a randomly chosen negative one [7]. The LDA classifier achieves an 84.6 % AUC score for the Training set and 83.3% for the Test set. These are very good scores for a dataset in its high dimensional nature.

¹True Positive is when a test result correctly indicates the presence of a condition. False Negative is a test result which wrongly indicates that a particular condition is absent.

KNN

K-nearest neighbours (KNN) is a discriminative approach which aims to model the conditional probability $p(y = k | x)$ directly. This method measures distances between each data point and the number of nearest neighbours. The pseudo algorithm for this method is found in Algorithm 1.

Algorithm 1 KNN algorithm [8]

Input: (X_i, C_i) where $i = 1, \dots, n$ be n data points, X_i denotes feature values and C_i denotes labels for each i . Integer k .

for $i=1, 2, \dots$ **do**
 | Calculate $d(x, x_i)$ where d is the Euclidean distance
end
Arrange d in increasing order
 Take first k distances from the sorted list
 Find those k -points corresponding to these k -distances
 Let k_i be the number of points in the i^{th} class among k points

for $i=1, 2, \dots$ **do**
 | **for** $j=1, 2, \dots$ **do**
 | **if** $k_i > k_j$ for all $i \neq j$ **then**
 | put x in class i
 | **end**
 | **end**
end
Return: Data points in different classes

To compute the KNN we use the `KNeighborsClassifier` function from `sklearn` which implements learning based on the k nearest neighbors of each query point. We choose the Euclidean distance as the metric since it provides the least error rate. The function uses uniform weights (all points in each neighbourhood are weighted equally). The algorithm used to compute the nearest neighbours was `ball_tree` with leaf size 30. This is a similar algorithm to 1, more about the algorithm can be found in [9]. The leaf size is 30 which is a good compromise between the computational time and amount of nodes created.

In order to tune the hyperparameter k , we try different k -values and for each value we calculate the mean accuracy rate (the times the predicted label matches the true label). To do this we again use cross validation with 5 folds as described in question 1. Using the function `GridSearchCV` we were able to plot the mean accuracy for each value of k for $k \in \{5, 10, 15, 20, 25, 30, 35, 40, 45, 50\}$ as seen in figure 9. We see that the accuracy stabilizes after $k = 25$ so we choose this value to avoid underfitting (decision boundary is closer to being linear as k increases). A smaller value would result in the decision boundary to be overly flexible and find patterns in the data that do not correspond to the Bayes decision boundary (overfitting).

We can assess the success of this classifier again by looking at the confusion matrix, error rate, ROC curve and AUC score. The confusion matrix can be seen in table 1. The error rate is 25%, 1% larger than the LDA classifier. The ROC curve seems to be slightly better than the LDA, with an AUC of 81.9% for the test set. For a model with only 12 nearest neighbours, these are very good statistics for the classifier.

Naive Bayes Classifier

The Naive Bayes classifier is a generative approach and it assumes all measured variables and features are independent given the label. The classifier makes the assumption that the features are conditionally independent for the class label:

$$p(x | y = k) = \prod_{i=1}^d p(x_i | y = k, \beta_{ik}) \quad (12)$$




Figure 9: Mean Accuracy against number of nearest neighbours for KNN model.

Since in this dataset we have real-valued features we decided to use the Gaussian distribution, ie:

$$p(x | y = k', \theta) = \prod_{j=1}^d \Phi(x_j | \mu_{jk}, \sigma_{jk}^2) \quad (13)$$

where Φ is the probability density function of the normal distribution.

We then use the Maximum Likelihood approach for parameter estimation, with the probability for a single data point x^i, y^i given by:

$$p(x^{(i)}, y^{(i)} | \theta, \pi) = p(y^{(i)} | \pi) \prod_{j=1}^d p(x_j^{(i)} | \theta_j) = \prod_{k=1}^K \pi_k^{1(y^{(i)}=k)} \prod_{j=1}^d \prod_{k=1}^K p(x_j^{(i)} | \theta_{jk})^{1(y^{(i)}=k)} \quad (14)$$

where π_k is the prior probability that $y^{(i)} = k$.

To fit the Naive Bayes classifier we will use the **GaussianNB** function from **sklearn** which implements the Gaussian Naive Bayes algorithm mentioned above. Fitting on the training set and testing on the test set we observe that the confusion matrix is much worse than the other two classifiers, with an error rate of 33%. This is mostly due to the number of instances the true label was 0 (before 1980) and was misclassified by 1 (after 1980). Looking at the ROC curve in figure 10 we see a slight decrease in the peak with comparison of the other curves. The AUC score for the test set was just at 75.6%, not bad in general but worse than our other two classifiers.

Conclusion

All three classifiers are able to classify songs before and after 1980 fairly good. By the nature of the problem, there is no hard decision boundary at 1980 that differentiates music in a fundamental way. That is to say, we can expect that songs of 1982 for example do not differ a lot with 1978. This suggests that an extension to the study would be to create softer margins, and group maybe decades together to see clear difference between the classes. Nevertheless, since there is a clear difference of the music style for songs of the 40s, 50s, 60s and 70s with everything after the 90s, we achieved very good classifiers that can predict the age of a song. Right-censoring the data so that we have roughly equal number of songs before and after the decision boundary really helped with the classification problem and meant that we can use the underlying assumptions of LDA and Naive Bayes. When trying to decide between the three classifiers we can refer to figure 11 to test that LDA seems to have the best ROC Curve with the highest AUC followed by KNN and then by Naive Bayes.




Figure 10: ROC curves of LDA (left), KNN (right) and Naive Bayes (right) for both Test and Training sets (30-70 % split)

| Classifier | LDA | KNN | NB |
|------------------|------|------|------|
| Computation time | 0.5s | 5.7s | 0.4s |

Table 2: Computational time of the fitting of all three classifiers for the same training subset of the data. The measurements were made on the same computer, with minimal background processes at the fitting procedure.

KNN was the only classifier which we could tune hyperparameters, which suggest that the other two models cannot be improved to a large extent, but we also believe we have found a good hyperparameter k for the KNN that does not overfit our data. Comparing the computational times of the three classifiers in table 2 it is clear that the KNN is the slowest and LDA and Naive Bayes to be almost identical with Naive Bayes having the edge. This makes sense since Naive Bayes makes the (crude) assumption of independent features given a label. We did not challenge this assumption because of how many features we had, but computing the correlation matrix did not show any clear correlation with the labels.

The KNN iteratively goes through every point and assigns it to a class which is computationally extensive. The ball tree algorithm does not work very well with high dimensions since the split of the tree at each node is based on all of the labels of our dataset. LDA works very well with high dimensional data.

In conclusion, the LDA seems to be the best classifier, satisfying both accuracy, computational cost, and clear interpretability of results.




Figure 11: ROC curves of Test sets of all 3 classifiers (LDA, KNN, Naive Bayes)

References

- [1] Dr Sarah Filippi. Machine learning lecture notes. *Imperial College London*, pages 30–40, October 2021.
- [2] Richard D Morey, Rink Hoekstra, Jeffrey N Rouder, Michael D Lee, and Eric-Jan Wagenmakers. The fallacy of placing confidence in confidence intervals. *Psychonomic bulletin & review*, 23(1):103–123, 02 2016. doi: 10.3758/s13423-015-0947-8. URL <https://pubmed.ncbi.nlm.nih.gov/26450628>.
- [3] Dheeru Dua and Casey Graff. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>.
- [4] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [5] Wikipedia. Confusion matrix — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Confusion%20matrix&oldid=1058352752>, 2022. [Online; accessed 29-January-2022].
- [6] Wikipedia. Receiver operating characteristic — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Receiver%20operating%20characteristic&oldid=1067917334>, 2022. [Online; accessed 29-January-2022].
- [7] Fawcett, Tom. An introduction to roc analysis. 2006. URL <https://www.math.ucdavis.edu/~saito/data/roc/fawcett-roc.pdf>.
- [8] Rahul Saxena. Knn classifier, introduction to k-nearest neighbor algorithm. <https://dataaspirant.com/k-nearest-neighbor-classifier-intro/>, 2022. [Online; accessed 29-January-2022].
- [9] Stephen M. Omohundro. Five balltree construction algorithms. Technical report, 1989.

A Code for Question 1

```
1 # %%
2 import pandas as pd
3 import numpy as np
4 import os
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 # from sklearn.preprocessing import PolynomialFeatures
8 # from sklearn.linear_model import LinearRegression
9 from sklearn.metrics import mean_squared_error
10 from sklearn.pipeline import make_pipeline
11 from sklearn.model_selection import GridSearchCV, KFold
12 # from sklearn.preprocessing import LabelEncoder # binary encoding
13 from sklearn.model_selection import train_test_split
14 # from sklearn.preprocessing import StandardScaler
15 # from sklearn.cluster import KMeans
16 import sklearn.pipeline
17 from sklearn.linear_model import Ridge
18 from sklearn.linear_model import RidgeCV
19 from scipy.stats import norm
20 seed = 222
21 np.random.seed(seed)
22
23 plt.rcParams['figure.figsize'] = (12,5)
24 plt.rcParams['figure.dpi'] = 80
```

```

25
26
27 # %%
28 d = pd.read_csv('data.csv')
29 d.describe()
30
31 # %%
32 sns.scatterplot(x='x', y='y', data = d, hue= 'z', s=20)
33 plt.savefig('Plots/my_data.pdf')
34 plt.show()
35
36 # %%
37 def create_design_matrix(x,z):
38
39     X = np.vstack((z, np.cos(x), z * np.cos(x), x, z * x, x * x,
40                    z * x * x, x * x * x, z * x * x * x ))
41     X = X.T
42     return(X)
43
44 X =create_design_matrix(d.x, d.z)
45 y = d.y
46
47
48 # %%
49 def get_coefs(m):
50     if (isinstance(m, sklearn.pipeline.Pipeline)):
51         m = m.steps[-1][1]
52     if m.intercept_ is None:
53         return m.coef_
54     return np.concatenate([[m.intercept_], m.coef_])
55
56
57 # %%
58 def model_fit(m, X, y, plot = False):
59
60     y_hat = m.predict(X)
61     rmse = mean_squared_error(y, y_hat, squared=False)
62     res = pd.DataFrame(
63         data = {'y': y, 'y_hat': y_hat, 'resid': y - y_hat}
64     )
65     if plot:
66         plt.figure(figsize=(12, 6))
67         plt.subplot(121)
68         sns.lineplot(x='y', y='y_hat', color="grey", data = pd.DataFrame(data
69         sns.scatterplot(x='y', y='y_hat', data=res).set_title("Fit plot")
70         plt.subplot(122)
71         sns.scatterplot(x='y', y='resid', data=res).set_title("Residual plot")
72         plt.suptitle("Model rmse = " + str(round(rmse, 4)), fontsize=16)
73         plt.savefig('Plots/model_summary.pdf')
74         plt.show()
75     return(rmse)
76
77
78 # %%
79 r = make_pipeline(
80     Ridge(alpha=1)
81 ).fit(X, y)
82
83 # %%
84 model_fit(r, X, y, plot=True)
85
86 # %%
87 alphas = np.linspace(0.01, 5, num=150)
88

```

```

89 gs = GridSearchCV(
90     make_pipeline(
91         Ridge()
92     ),
93     param_grid={'ridge__alpha': alphas},
94     cv=KFold(5, shuffle=True, random_state=1234),
95     scoring="neg_root_mean_squared_error"
96 ).fit(X, y)
97
98 # %%
99 print(gs.best_params_)
100 model_fit(gs.best_estimator_, X, y, plot=True)
101 best_alpha = gs.best_params_['ridge__alpha']
102
103 # %%
104 cv_res = pd.DataFrame(
105     data = gs.cv_results_
106 ).filter(
107     regex = '(split[0-9]+|mean)_test_score'
108 ).assign(
109     # Add the alphas as a column
110     alpha = alphas
111 )
112 cv_res.update(
113     # Convert negative rmse to positive
114     -1 * cv_res.filter(regex = '_test_score')
115 )
116 sns.lineplot(x='alpha', y='mean_test_score', data=cv_res)
117 plt.xlabel('$\lambda$')
118 plt.ylabel('Negative RMSE')
119 # plt.savefig('Plots/single_1.pdf')
120 plt.show()
121
122 # %%
123 d1 = cv_res.melt(
124     id_vars=('alpha', 'mean_test_score'),
125     var_name='fold',
126     value_name='rmse'
127 )
128
129 sns.lineplot(x='alpha', y='rmse', color='black', ci = None, data = d1) # Plot
130 sns.lineplot(x='alpha', y='rmse', hue='fold', data = d1) # Plot the curves for
131 plt.xlabel('$\lambda$')
132 # plt.savefig('Plots/splits.pdf')
133 plt.show()
134
135 # %%
136 m = make_pipeline(
137     Ridge(alpha=best_alpha, fit_intercept = False, solver='svd' )
138 ).fit(X, y)
139
140 # %%
141 model_fit(m, X, y, plot= True)
142
143 # %%
144 new_x = np.arange(-7, 7.1, step=0.1)
145
146 X_new1 = create_design_matrix(new_x, np.zeros(new_x.shape))
147 X_new2 = create_design_matrix(new_x, np.ones(new_x.shape))
148
149 # %%
150 y_hat1 = m.predict(X_new1)
151 y_hat2 = m.predict(X_new2)
152

```

```

153
154 # %%
155 plt.figure(figsize=(12, 6))
156 plt.subplot(121)
157 sns.scatterplot(x='x', y='y', data = d[d.z == 0], s=20)
158 plt.plot(new_x, y_hat1, color = 'red', label = 'z=0')
159 plt.legend()
160 plt.subplot(122)
161 sns.scatterplot(x='x', y='y', data = d[d.z == 1], s=20)
162 plt.plot(new_x, y_hat2, color = 'green', label='z=1')
163 plt.legend()
164 # plt.savefig('Plots/two_models_ridge.pdf')
165 plt.show()
166
167 # %% [markdown]
168 # ## Part 2
169
170 # %%
171 sigma_sq = 0.1
172
173 # %%
174 # see equations page 31 of slides
175 def solve_posterior(
176     X: np.ndarray, y: np.ndarray, sigma_sq: np.ndarray, alpha: np.ndarray
177 ) -> np.ndarray:
178     n = X.shape[0]
179     p = X.shape[1]
180     mu = np.linalg.solve(X.T @ X + sigma_sq / alpha * np.eye(p), X.T @ y)
181     Sigma = sigma_sq * np.linalg.solve(
182         X.T @ X + (sigma_sq / alpha) * np.eye(p), np.eye(p)
183     )
184     return mu, Sigma
185
186 # %%
187 alphas = np.linspace(0.001, 0.2, 150)
188 mus = np.empty((alphas.shape[0], 9))
189 for i, a in enumerate(alphas):
190     mus[i] = solve_posterior(X, y, sigma_sq, a)[0]
191 mus.shape
192
193
194 # %%
195 my_lst = [(0,0), (0,1), (0,2), (1,0), (1, 1), (1,2), (2,0), (2,1), (2,2)]
196 all_par = ['$a$', '$b_1$', '$b_2$', '$c_1$', '$c_2$', '$d_1$', '$d_2$', '$e_1$', '$e_2$']
197 fig, axs = plt.subplots(3, 3)
198 for i, (t, p) in enumerate(zip(my_lst, all_par)):
199     axs[t].plot(alphas, mus[:,i])
200     axs[t].set_title(p)
201
202 for ax in axs.flat:
203     ax.set(xlabel='alpha', ylabel='Posterior Mean')
204
205 for ax in axs.flat:
206     ax.label_outer()
207
208 # fig.savefig('Plots/all_par.pdf')
209
210 # %%
211 mu, sigma = solve_posterior(X, y, sigma_sq, 1)
212
213 # %%
214 pred_1 = X_new1 @ mu
215 pred_2 = X_new2 @ mu
216 pred_1_var = X_new1 @ sigma @ X_new1.T + sigma_sq

```

```

217 pred_1_var = X_new2 @ sigma @ X_new2.T + sigma_sq
218
219 # %%
220 plt.scatter(d.x,d.y)
221 plt.plot(new_x, pred_1, color = 'red')
222 plt.plot(new_x, pred_2, color = 'green')
223
224
225 # %%
226 cred_int_1 = np.empty((X_new1.shape[0],2))
227 for i in range(0, X_new1.shape[0]):
228     big_sigma = X_new1[i] @ sigma @ X_new1[i].T + sigma_sq
229     cred_int_1[i,0] = norm.ppf(0.025, pred_1[i], big_sigma)
230     cred_int_1[i,1] = norm.ppf(0.975, pred_1[i], big_sigma)
231
232 cred_int_2 = np.empty((X_new2.shape[0],2))
233 for i in range(0, X_new2.shape[0]):
234     big_sigma = X_new2[i] @ sigma @ X_new2[i].T + sigma_sq
235     cred_int_2[i,0] = norm.ppf(0.025, pred_2[i], big_sigma)
236     cred_int_2[i,1] = norm.ppf(0.975, pred_2[i], big_sigma)
237
238 # %%
239 plt.figure(figsize=(12, 6))
240 plt.subplot(121)
241 sns.scatterplot(x='x', y='y', data = d[d.z == 0], s=20)
242 plt.plot(new_x, pred_1, color = 'red', linewidth=0.5, label='z=0')
243 plt.fill_between(new_x, cred_int_1[:,0], cred_int_1[:,1], alpha=0.3, color='blue')
244 plt.legend()
245 plt.subplot(122)
246 sns.scatterplot(x='x', y='y', data = d[d.z == 1], s=20)
247 plt.plot(new_x, pred_2, color = 'green', linewidth=0.5, label='z=1')
248 plt.fill_between(new_x, cred_int_2[:,0], cred_int_2[:,1], alpha=0.3, color='blue')
249 plt.legend()
250 # plt.savefig('Plots/two_models_bayes.pdf')
251 plt.show()
252
253 # %%
254 X_trial = create_design_matrix(np.array([-5,7, -5, 7]), np.array([0,1, 1,0]))
255
256 # %%
257 pred_bayesian1 = X_trial @ mu
258
259 pred_ridge = m.predict(X_trial)
260
261 # %%
262 ridge_alpha = sigma_sq / best_alpha
263 ridge_alpha
264
265 # %%
266 def calc_all(i):
267     new_mu = solve_posterior(X,y,sigma_sq,0.001)[0]
268     new_sigma = solve_posterior(X,y,sigma_sq,0.001)[1]
269     mean_small = X_trial[i] @ new_mu
270     sigma_small = X_trial[i] @ new_sigma @ X_trial[i].T
271     new_mu = solve_posterior(X,y,sigma_sq,1)[0]
272     new_sigma = solve_posterior(X,y,sigma_sq,1)[1]
273     mean_medium = X_trial[i] @ new_mu
274     sigma_medium = X_trial[i] @ new_sigma @ X_trial[i].T
275     new_mu = solve_posterior(X,y,sigma_sq,ridge_alpha)[0]
276     new_sigma = solve_posterior(X,y,sigma_sq,ridge_alpha)[1]
277     mean_ridge = X_trial[i] @ new_mu
278     sigma_ridge = X_trial[i] @ new_sigma @ X_trial[i].T
279     return(mean_small, sigma_small, mean_medium, sigma_medium, mean_ridge, sigma_ridge)
280

```

```

281 # %%
282 plt.subplots_adjust(hspace=0.5)
283
284 plt.subplot(221)
285 mean_small, sigma_small, mean_medium, sigma_medium, mean_ridge, sigma_ridge = c
286 t1 = np.linspace(mean_small - 15*sigma_small, mean_small + 15*sigma_small, 100)
287 plt.plot(t1, norm.pdf(t1, mean_small, sigma_small),label='alpha: 0.001' )
288 plt.plot(t1, norm.pdf(t1, mean_medium, sigma_medium), label='alpha: 1')
289 plt.plot(t1, norm.pdf(t1, mean_ridge, sigma_ridge),label='alpha: 0.0822' )
290 plt.axvline(x=pred_bayesian1[0], color='red', label = 'Ridge Estimate')
291 plt.xlabel('$y^{*}$')
292 plt.title('X: -5, Z: 0')
293 # plt.legend()
294
295 plt.subplot(222)
296 mean_small, sigma_small, mean_medium, sigma_medium, mean_ridge, sigma_ridge = c
297 t1 = np.linspace(mean_small - 15*sigma_small, mean_small + 15*sigma_small, 100)
298 plt.plot(t1, norm.pdf(t1, mean_small, sigma_small),label='alpha: 0.001' )
299 plt.plot(t1, norm.pdf(t1, mean_medium, sigma_medium), label='alpha: 1')
300 plt.plot(t1, norm.pdf(t1, mean_ridge, sigma_ridge),label='alpha: 0.0822' )
301 plt.axvline(x=pred_bayesian1[1], color='red', label = 'Ridge Estimate')
302 plt.title('X: 7, Z: 1')
303 plt.xlabel('$y^{*}$')
304
305 plt.legend()
306
307 plt.subplot(223)
308 mean_small, sigma_small, mean_medium, sigma_medium, mean_ridge, sigma_ridge = c
309 t1 = np.linspace(mean_small - 15*sigma_small, mean_small + 55*sigma_small, 100)
310 plt.plot(t1, norm.pdf(t1, mean_small, sigma_small),label='alpha: 0.001' )
311 plt.plot(t1, norm.pdf(t1, mean_medium, sigma_medium), label='alpha: 1')
312 plt.plot(t1, norm.pdf(t1, mean_ridge, sigma_ridge),label='alpha: 0.0822' )
313 plt.axvline(x=pred_bayesian1[2], color='red', label = 'Ridge Estimate')
314 plt.xlabel('$y^{*}$')
315
316 plt.title('X: -5, Z: 1')
317 # plt.legend()
318
319 plt.subplot(224)
320 mean_small, sigma_small, mean_medium, sigma_medium, mean_ridge, sigma_ridge = c
321 t1 = np.linspace(mean_small - 15*sigma_small, mean_small + 15*sigma_small, 100)
322 plt.plot(t1, norm.pdf(t1, mean_small, sigma_small),label='alpha: 0.001' )
323 plt.plot(t1, norm.pdf(t1, mean_medium, sigma_medium), label='alpha: 1')
324 plt.plot(t1, norm.pdf(t1, mean_ridge, sigma_ridge),label='alpha: 0.0822' )
325 plt.axvline(x=pred_bayesian1[3], color='red', label = 'Ridge Estimate')
326 plt.title('X: 7, Z: 0')
327 plt.xlabel('$y^{*}$')
328 # plt.legend()
329
330 # plt.savefig('Plots/post_all.pdf')
331
332
333 # %%

```

B Code for Question 2

```

1 # %%
2 import pandas as pd
3 import numpy as np
4 import os
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.metrics import mean_squared_error
8 from sklearn.pipeline import make_pipeline

```



```

9  from sklearn.model_selection import KFold
10 from sklearn.model_selection import train_test_split
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.metrics import confusion_matrix
13 from sklearn.metrics import roc_curve, precision_recall_curve
14 from sklearn.metrics import roc_auc_score
15 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
16 from sklearn.naive_bayes import GaussianNB
17 from sklearn.neighbors import KNeighborsClassifier
18 from sklearn.model_selection import GridSearchCV, KFold
19
20 np.random.seed(222)
21
22 # %%
23 plt.rcParams['figure.figsize'] = (12,5)
24 plt.rcParams['figure.dpi'] = 80
25
26 # %%
27 d = pd.read_csv("dataQ2.csv")
28
29 # %%
30 df = d.groupby('V1').count()
31
32 clr = ['grey' if (x < 1980) else '#3282bd' for x in df.index ]
33
34 plt.figure(figsize=(20,10))
35 plt.xticks(rotation=90)
36 # plot = sns.barplot(x=df.index, y = df.V3, color='#3282bd')
37 plot = sns.barplot(x=df.index, y = df.V3, palette=clr)
38 plot.set_ylabel('Counts', fontsize=20, labelpad=20)
39 plot.set_xlabel('Year', fontsize=20, labelpad=20)
40 plt.savefig('Plots/song_year.pdf')
41 plt.show()
42
43 # %%
44 d['V1'] = np.where(d['V1'] >= 1980, 1,0)
45 d
46
47 # %% [markdown]
48 # ### LDA
49
50 # %%
51 y = d['V1']
52 X = d.loc[:,d.columns != 'V1']
53
54
55
56 # %%
57 TEST_SIZE = 0.3
58
59 # train/test split
60 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=TEST_SIZE)
61
62 # %%
63 sc = StandardScaler()
64 X_train = sc.fit_transform(X_train)
65 X_test = sc.transform(X_test)
66
67 # %%
68 lda = LDA(solver='svd', tol=1.e-4)
69 song_lda = lda.fit(X=X_train, y=y_train)
70
71 # %%
72 print(f"Prior probabilities of groups: \n {song_lda.priors_} \n")

```

```

73
74 # %%
75 lda_train = song_lda.predict(X_train)
76 lda_test = song_lda.predict(X_test)
77
78 # %%
79 cf_train = confusion_matrix(y_train, lda_train)
80 pd.DataFrame(cf_train, columns=["FALSE", "TRUE"], index=["FALSE", "TRUE"])
81
82 # %%
83 cf_test = confusion_matrix(y_test, lda_test)
84 pd.DataFrame(cf_test, columns=["FALSE", "TRUE"], index=["FALSE", "TRUE"])
85
86 # %%
87 from sklearn.metrics import roc_curve
88 fpr_1,tpr_1,thresh=roc_curve(y_test,lda.predict_proba(X_test)[:,-1])
89 auc_score=roc_auc_score(y_test,lda.predict_proba(X_test)[:,-1])
90
91 fpr_12,tpr_12,thresh2=roc_curve(y_train,lda.predict_proba(X_train)[:,-1])
92 auc_score2=roc_auc_score(y_train,lda.predict_proba(X_train)[:,-1])
93
94 print(f'AUC of Train: {auc_score2} \nAUC of Test: {auc_score}')
95
96 plt.plot(fpr_1,tpr_1, label='Test')
97 plt.plot(fpr_12, tpr_12, label='Train')
98 plt.xlabel('fpr1')
99 plt.ylabel('TPR')
100 plt.legend()
101
102 # %% [markdown]
103 # ## KNN
104
105 # %%
106 d2 = d.copy()
107 y = d2['V1']
108 X = d2.loc[:,d.columns != 'V1']
109
110 TEST_SIZE = 0.3
111
112 # train/test split
113 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=TEST_SIZE,
114
115 sc = StandardScaler()
116 X_train = sc.fit_transform(X_train)
117 X_test = sc.transform(X_test)
118
119
120 # X = sc.fit_transform(X)
121
122 # %%
123 KNeighborsClassifier().get_params().keys()
124
125 # %%
126 #trial:
127 ks = np.arange(5, 35)
128 ks = np.array([5,10,15,20,25,30,35,40,45,50])
129 gs = GridSearchCV(
130     make_pipeline(
131         StandardScaler(),
132         KNeighborsClassifier()
133     ),
134     param_grid={'kneighborsclassifier__n_neighbors': ks},
135     cv=KFold(5, shuffle=True, random_state=1234)
136 ).fit(X, y)

```

```

137
138
139 # %%
140 print(gs.best_params_)
141 test_score = gs.cv_results_['mean_test_score']
142
143 plt.figure(figsize=(12, 6))
144 plt.plot(ks, test_score, color='#3282bd', linestyle='dashed', marker='o',
145         markerfacecolor='black', markersize=10)
146 plt.xlabel('k Value')
147 plt.ylabel('Mean Accuracy')
148 plt.savefig('Plots/errorknn.pdf')
149 plt.show()
150
151 # %%
152 classifier = KNeighborsClassifier(n_neighbors=25)
153 classifier.fit(X_train, y_train)
154
155 y_pred = classifier.predict(X_test)
156
157
158 # %%
159 cf_test = confusion_matrix(y_test, y_pred)
160 pd.DataFrame(cf_test, columns=["FALSE", "TRUE"], index=["FALSE", "TRUE"])
161
162 # %%
163 (479+459)/(1478+459+479+1341)
164
165 # %%
166 fpr_2, tpr_2, thresh = roc_curve(y_test, classifier.predict_proba(X_test)[:,-1])
167 auc_score = roc_auc_score(y_test, classifier.predict_proba(X_test)[:,-1])
168
169 fpr_22, tpr_22, thresh2 = roc_curve(y_train, classifier.predict_proba(X_train)[:,-1])
170 auc_score2 = roc_auc_score(y_train, classifier.predict_proba(X_train)[:,-1])
171
172 print(f'AUC of Train: {auc_score2} \nAUC of Test: {auc_score}')
173
174 plt.plot(fpr_2, tpr_2, label='Test')
175 plt.plot(fpr_22, tpr_22, label='Train')
176 plt.legend()
177
178 # %% [markdown]
179 # ## Naive Bayes
180
181 # %%
182 d3 = d.copy()
183
184 y = d3['V1']
185 X = d3.loc[:, d3.columns != 'V1']
186
187 TEST_SIZE = 0.3 # smaller training set to reduce computational cost
188
189 # train/test split
190 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=TEST_SIZE)
191
192 sc = StandardScaler()
193 X_train = sc.fit_transform(X_train)
194 X_test = sc.transform(X_test)
195
196 # %%
197 from sklearn.naive_bayes import *
198
199 NB = GaussianNB()
200 # NB = MultinomialNB()

```

```

201 # NB = ComplementNB()
202
203 nb_fit = NB.fit(X_train, y_train)
204
205 y_pred = nb_fit.predict(X_test)
206
207 # %%
208 print("Number of mislabeled points out of a total %d points : %d" % (X_test.shape[0],
209
210 # %%
211 cf_test = confusion_matrix(y_test, y_pred)
212 pd.DataFrame(cf_test, columns=["FALSE", "TRUE"], index=["FALSE", "TRUE"]))
213
214 # %%
215 fpr_3, tpr_3, thresh = roc_curve(y_test, NB.predict_proba(X_test)[:, -1])
216 auc_score = roc_auc_score(y_test, NB.predict_proba(X_test)[:, -1])
217
218 fpr_32, tpr_32, thresh2 = roc_curve(y_train, NB.predict_proba(X_train)[:, -1])
219 auc_score2 = roc_auc_score(y_train, NB.predict_proba(X_train)[:, -1])
220
221 print(f'AUC of Train: {auc_score2} \nAUC of Test: {auc_score}')
222
223 plt.plot(fpr_3, tpr_3, label='Test')
224 plt.plot(fpr_32, tpr_32, label='Train')
225 plt.legend()
226
227 # %%
228 nb_fit.score(X_test, y_test)
229
230 # %%
231 t = np.linspace(0, 1, 50)
232 plt.plot(fpr_1, tpr_1, label='LDA')
233 # plt.plot(fpr_12, tpr_12, label='Train')
234 plt.plot(fpr_2, tpr_2, label='KNN')
235 # plt.plot(fpr_22, tpr_22, label='Train')
236 plt.plot(fpr_3, tpr_3, label='Naive Bayes')
237 # plt.plot(fpr_32, tpr_32, label='Train')
238 plt.plot(t, t, '--', alpha=0.3, color='grey')
239 plt.legend()
240 plt.xlabel('FPR')
241 plt.ylabel('TPR')
242 plt.savefig('Plots/together_plots.pdf')
243 plt.show()
244
245
246 # %%
247 t = np.linspace(0, 1, 50)
248 plt.subplot(131)
249 plt.plot(fpr_1, tpr_1, label='Test')
250 plt.plot(fpr_12, tpr_12, label='Train')
251 plt.plot(t, t, '--', alpha=0.3, color='grey')
252 plt.title('LDA')
253 plt.xlabel('FPR')
254 plt.ylabel('TPR')
255 plt.legend()
256 plt.subplot(132)
257 plt.plot(fpr_2, tpr_2, label='Test')
258 plt.plot(fpr_22, tpr_22, label='Train')
259 plt.plot(t, t, '--', alpha=0.3, color='grey')
260 plt.title('KNN')
261 plt.xlabel('FPR')
262 plt.legend()
263 plt.subplot(133)
264 plt.plot(fpr_3, tpr_3, label='Test')

```

```
265 plt.plot(fpr_32, tpr_32, label='Train')
266 plt.plot(t, t, '--', alpha=0.3, color='grey')
267 plt.title('Naive Bayes')
268 plt.xlabel('FPR')
269 plt.legend()
270 plt.savefig('Plots/allroc.pdf')
271 plt.show()
272
273 # %%
```